KubeCon | CloudNativeCon

North America 2018

Life of a Kubernetes Watch Event

Haowei Cai (roycaihw@github), Google
Wenjia Zhang (wenjiaswe@github), Google

About us

Haowei Cai(@roycaihw)

Software Engineer for Google Cloud. He is one of the owners of Kubernetes Python client library and an active Kubernetes SIG API Machinery contributor.

Wenjia Zhang (@wenjiaswe)

Software Engineer on Kubernetes team at Google. Active contributor for Kubernetes and etcd open source projects.

Agenda

- **What** is a Kubernetes Watch Event?
- **Why** is Watch Event important for Kubernetes?
- **How** is the life of a Kubernetes Watch Event?
- **Key Takeaways**

What is a Kubernetes Watch Event?

What is Watch?

Watch is an incremental change notification feed

Watch vs. Poll

Low latency

Single connection

Watch

Poll

Extra load
Extra latency

Multiple connection

Watch vs. Poll

Kubelet on nodes:

- Previous: periodically poll kube-apiserver for secrets and configmaps
- Now: watch individual secrets
- OSS PR: [Kubelet watches necessary secrets/configmaps instead of periodic polling #64752]

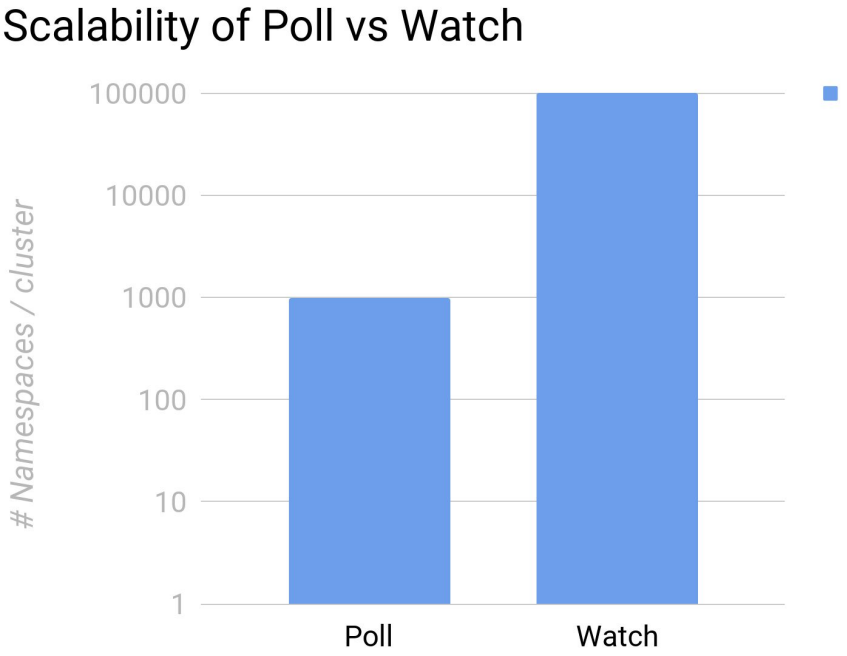

What is Event?

A Single event to a watched resource

```
// Event represents a single event to a watched resource.
// +k8s:deepcopy-gen=true
type Event struct {
   Type EventType

   // Object is:
   //  * If Type is Added or Modified: the new state of the object.
   //  * If Type is Deleted: the state of the object immediately before deletion.
   //  * If Type is Error: *api.Status is recommended; other types may make sense
   //    depending on context.
   Object runtime.Object
}
```

Watched resource: node — **Event:** add a node

Watched resource: pod — **Event**: modify a pod

Watched resource: replicaSet — **Event**: delete a replicaSet

Why is Watch Event important for Kubernetes?

Kubernetes core design concept

Level Triggering and Soft Reconciliation

Declarative configuration

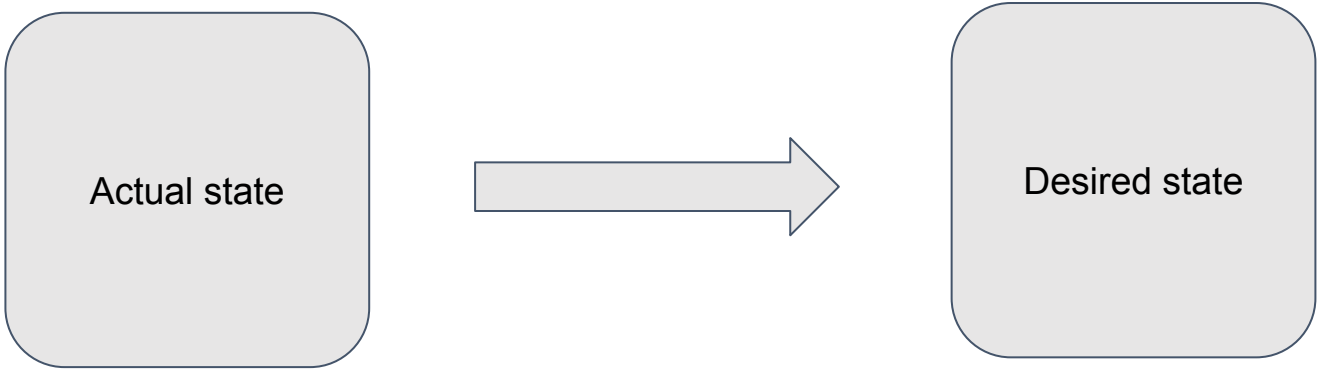

Declarative configuration

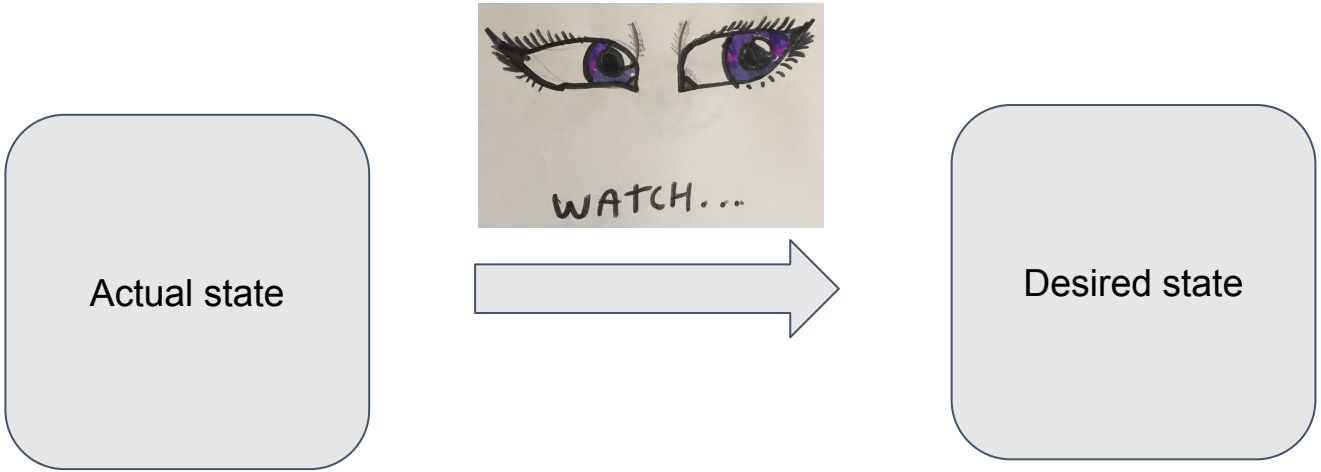

How is the lifecycle of a K8s watch event?

Fingerprint of events: resourceVersion

resourceVersion:

a string that identifies the common version of the objects returned by in a list. This value MUST be treated as opaque by clients and passed unmodified back to the server. A resource version is valid on a single kind of resource across namespaces.

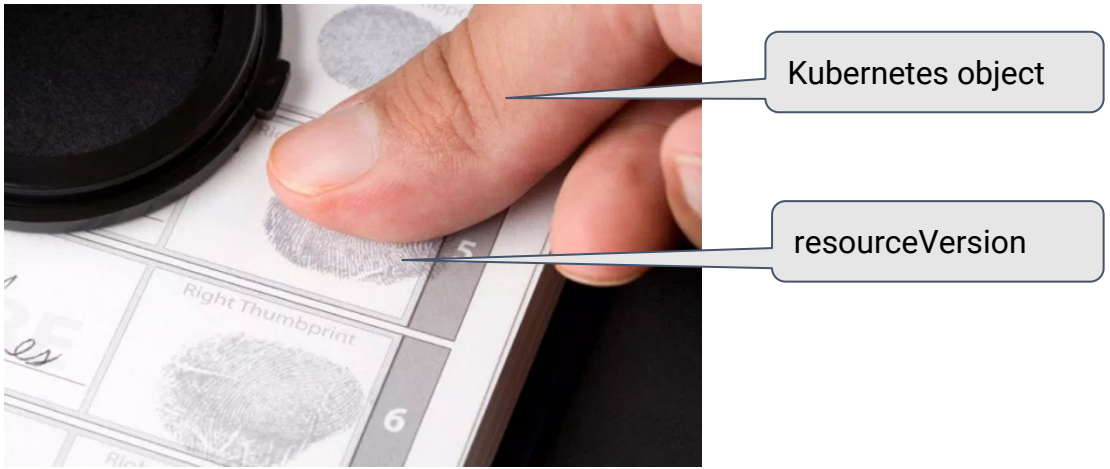

Lifecycle of a K8s watch event

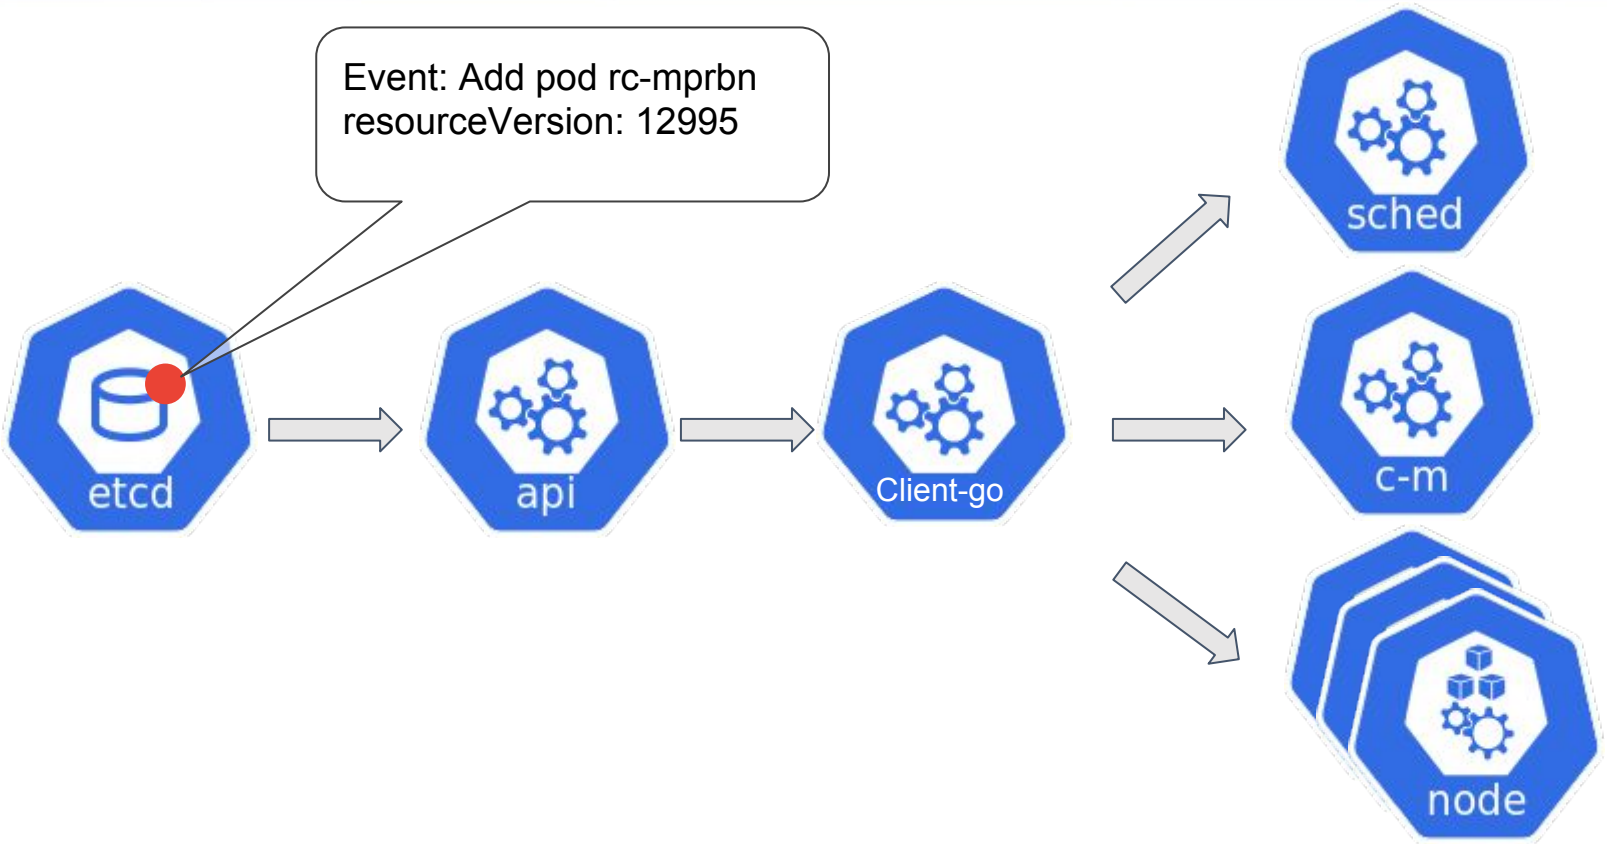

Lifecycle of a K8s watch event

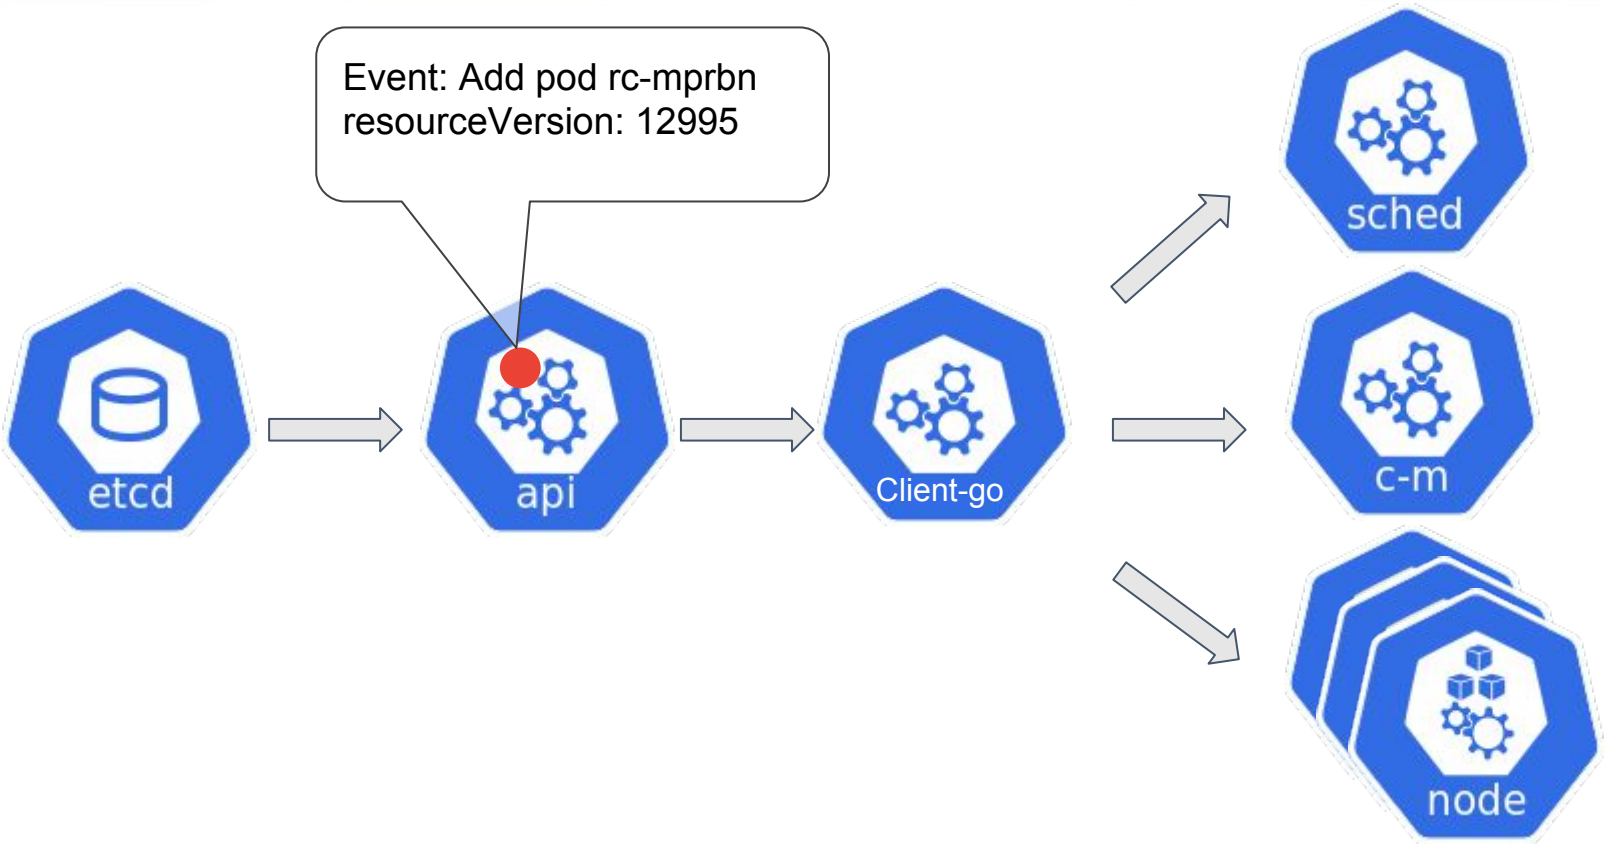

Lifecycle of a K8s watch event

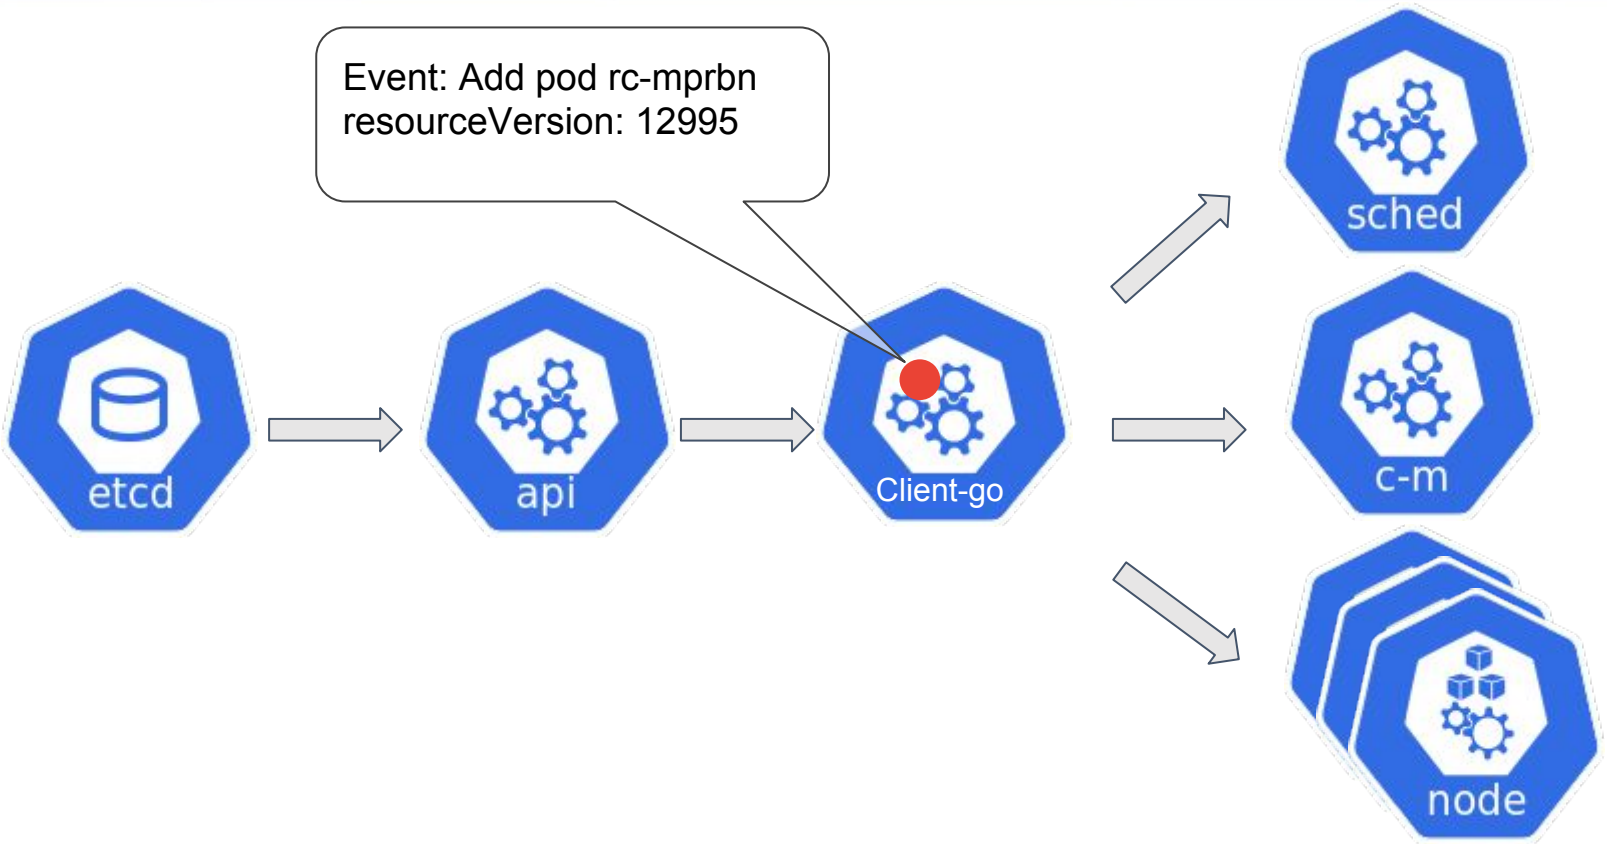

Lifecycle of a K8s watch event

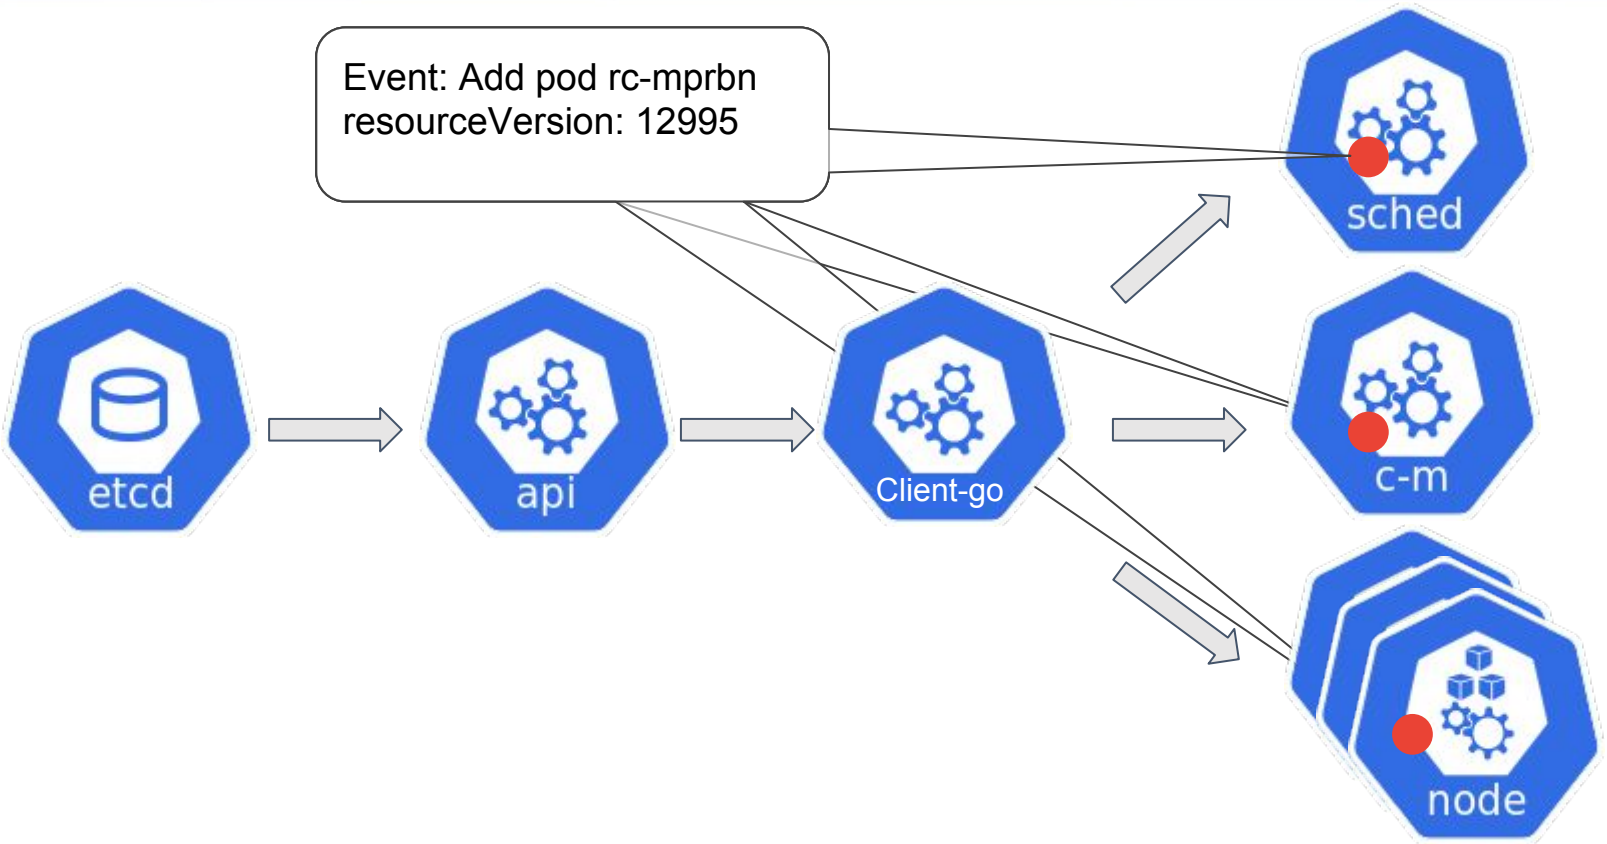

Watch Event on etcd

Watch in etcd

Etcd **watch** feature provides an event-based interface for asynchronously monitoring changes to keys.

```
tx01 $ etcdctl --endpoints=$ENDPOINTS
```

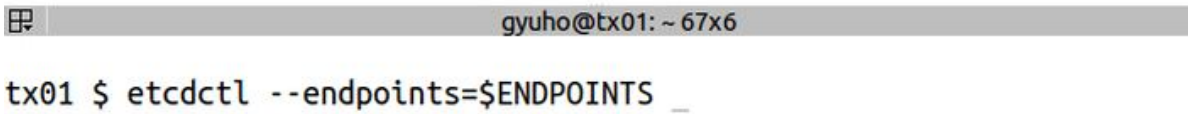

Revision

Revision (etcd) = resourceVersion (apiserver)

- The key space maintains multiple revisions.
- Each atomic mutative operation creates a new revision on the key space.
- All data held by previous revisions remains unchanged.
- If the store is compacted to save space, revisions before the compact revision will be removed.
- Revisions are monotonically increasing over the lifetime of a cluster.

*this is a current implementation detail of the etcd storage layer, and does not guarantee resourceVersion will be numeric or monotonically increasing in the future

Watch event on etcd

Watches make three guarantees about events:

- Ordered - events are ordered by revision
- Reliable - a sequence of events will never drop any subsequence of events
- Atomic - a list of events is guaranteed to encompass complete revisions

Watch Event on Kube Api-server

Watch Event on Kube Api-server

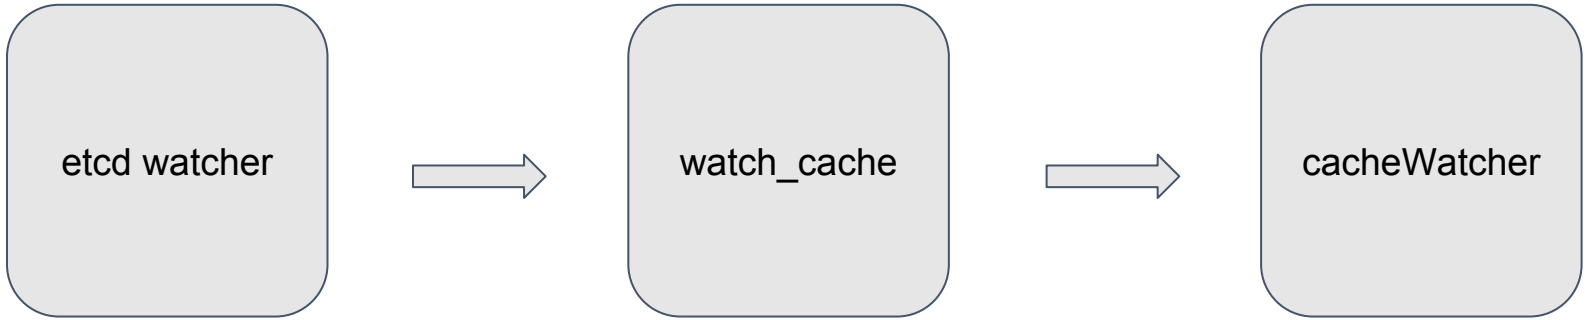

Watch Event on Kube Api-server

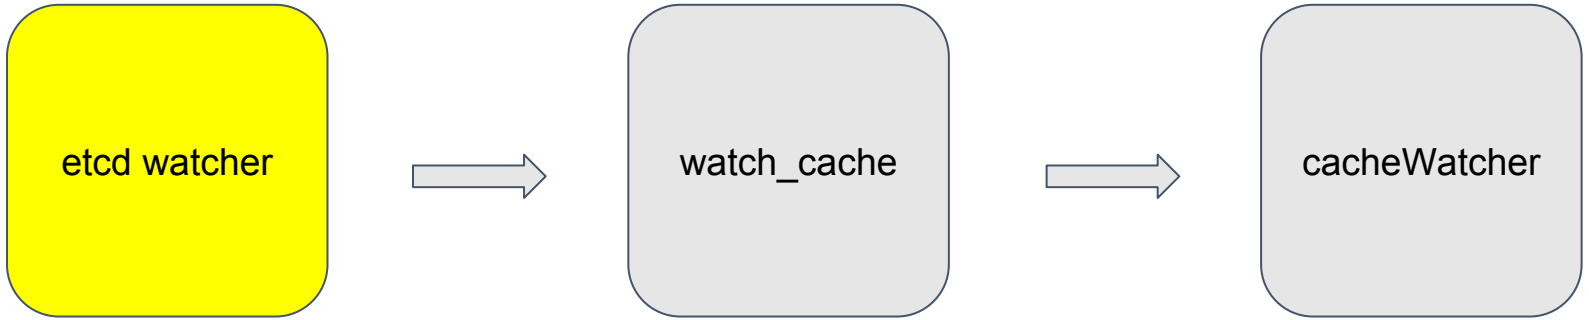

Watch Event on Kube Api-server

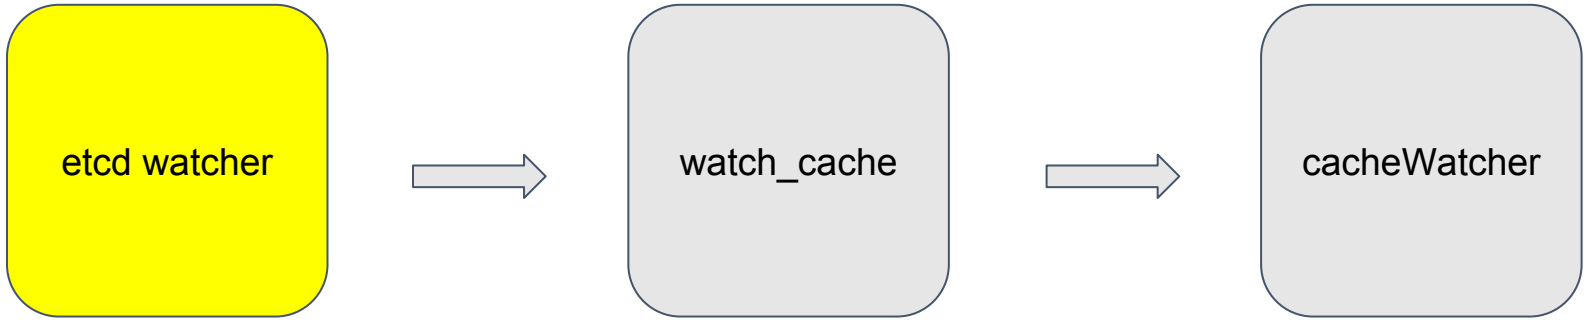

Watch Event on Kube Api-server

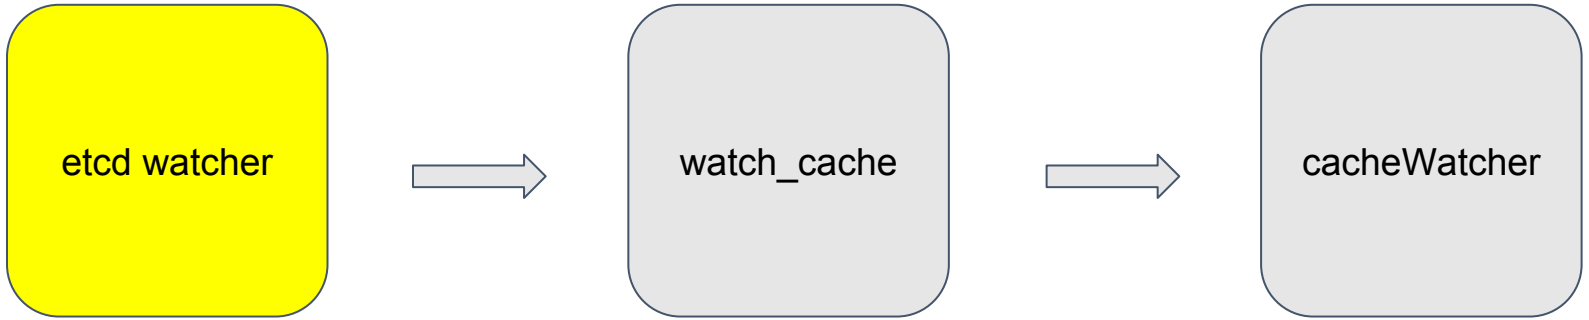

- Flow control…
- etcd3 Event -> APIserver Event...

Watch Event on Kube Api-server

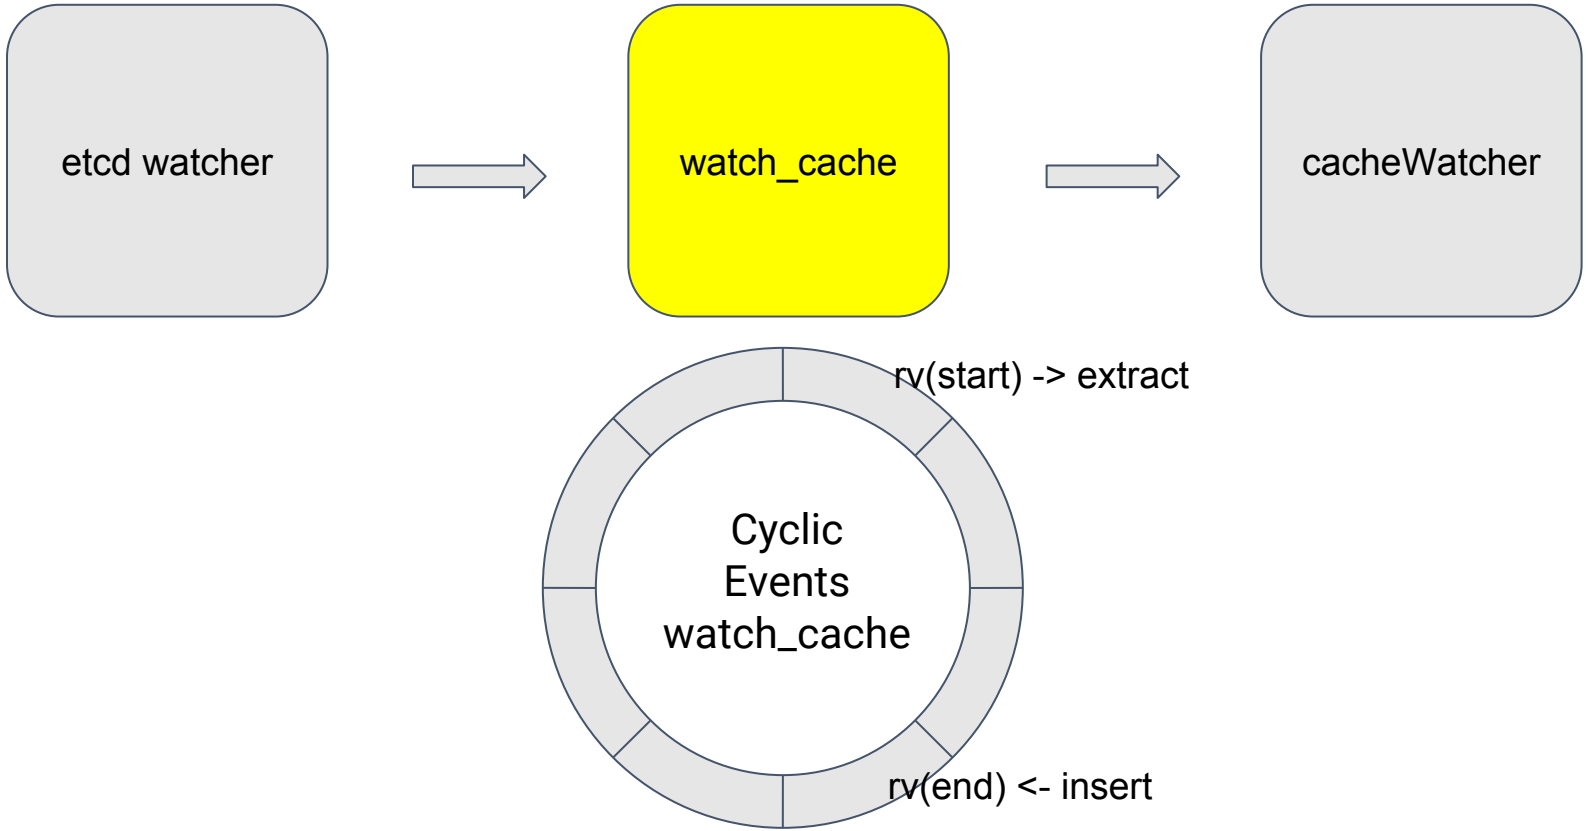

Watch Event on Kube Api-server

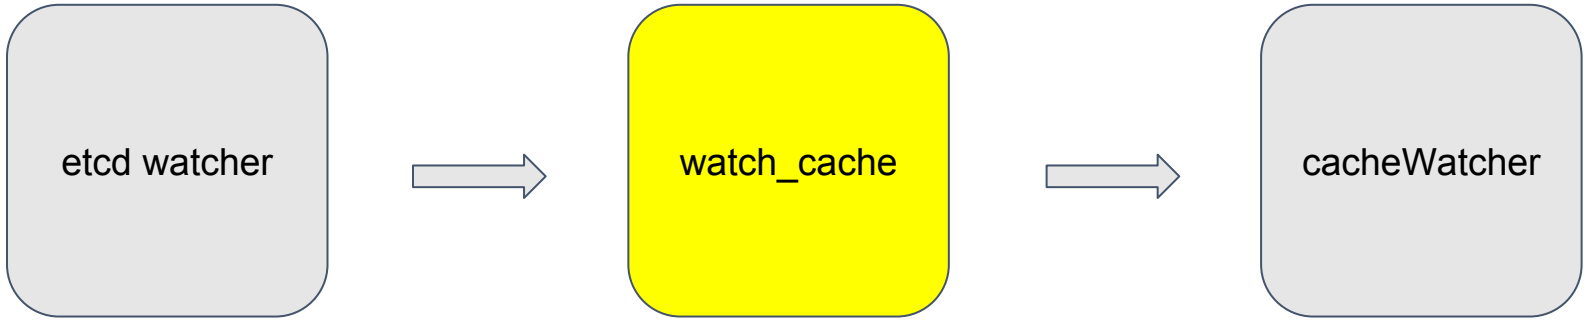

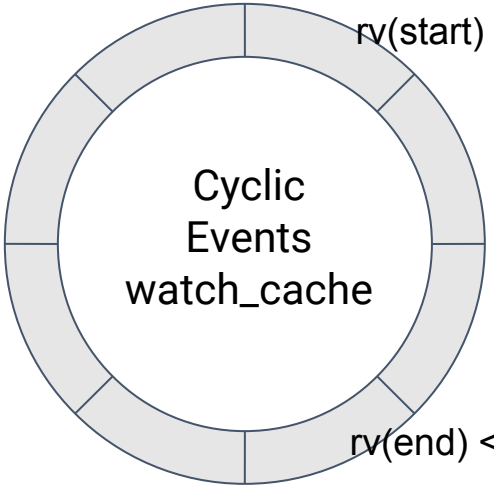

- watch_cache / resource type
- capacity / watch_cache
- List operation

Watch Event on Kube Api-server

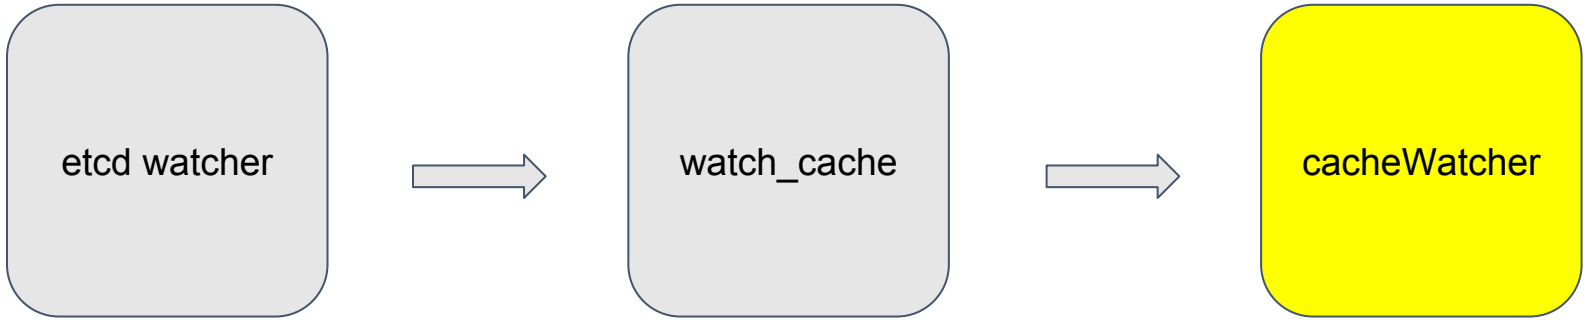

Serving Client

Watch Event in Client-go

What is Client-go?

Clientset
Dynamic Client
REST Client
Informer
...

https://github.com/kubernetes/client-go

- Go clients for talking to a kubernetes cluster
- Used by Kubernetes itself

What is Informer?

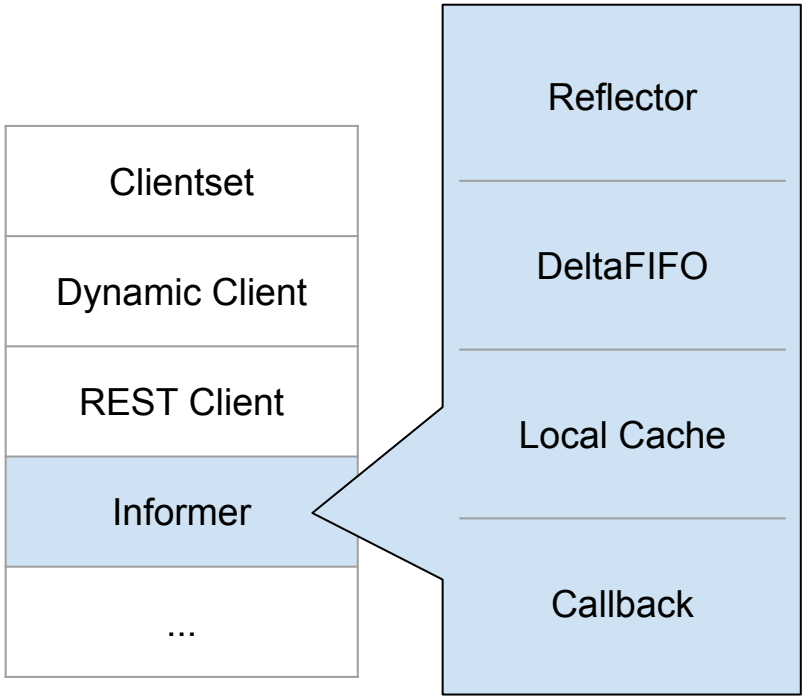

k8s.io/client-go/tools/cache
k8s.io/client-go/informers

- Useful component for building event-oriented controllers
- Used by control plane controllers, kubelet, etc.
- Reflector used by kube-apiserver watch cache

Kubernetes controller workflow

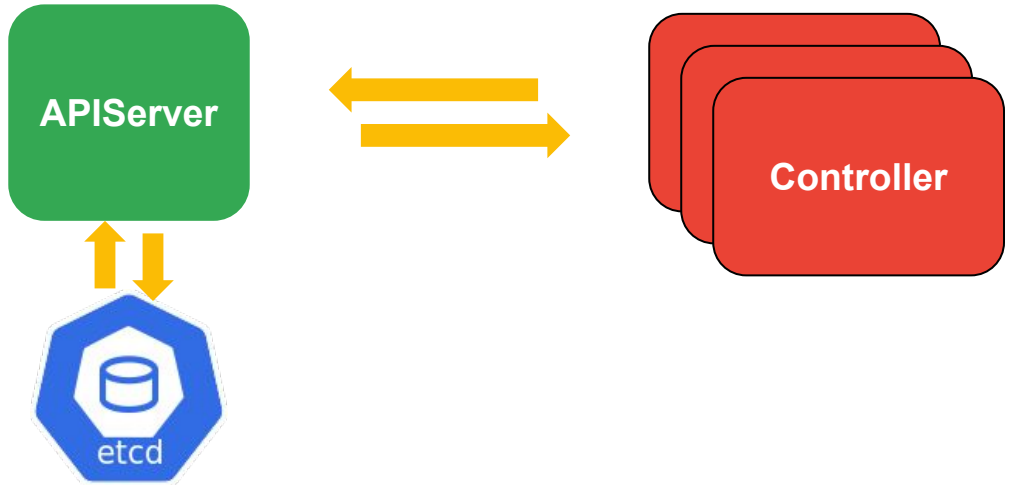

Kubernetes controller workflow

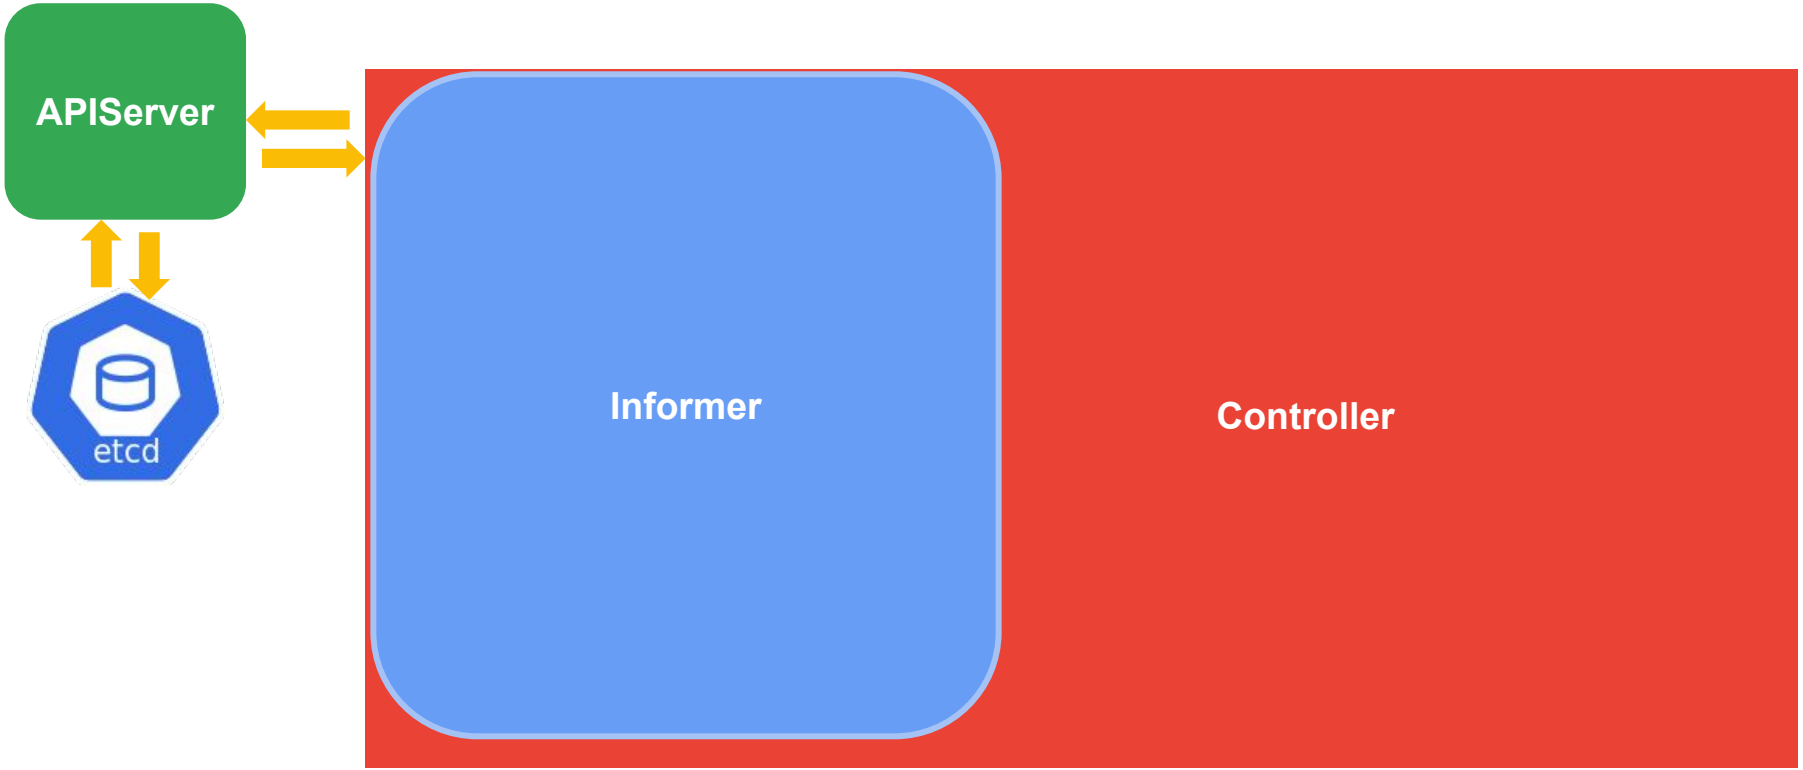

Kubernetes controller workflow

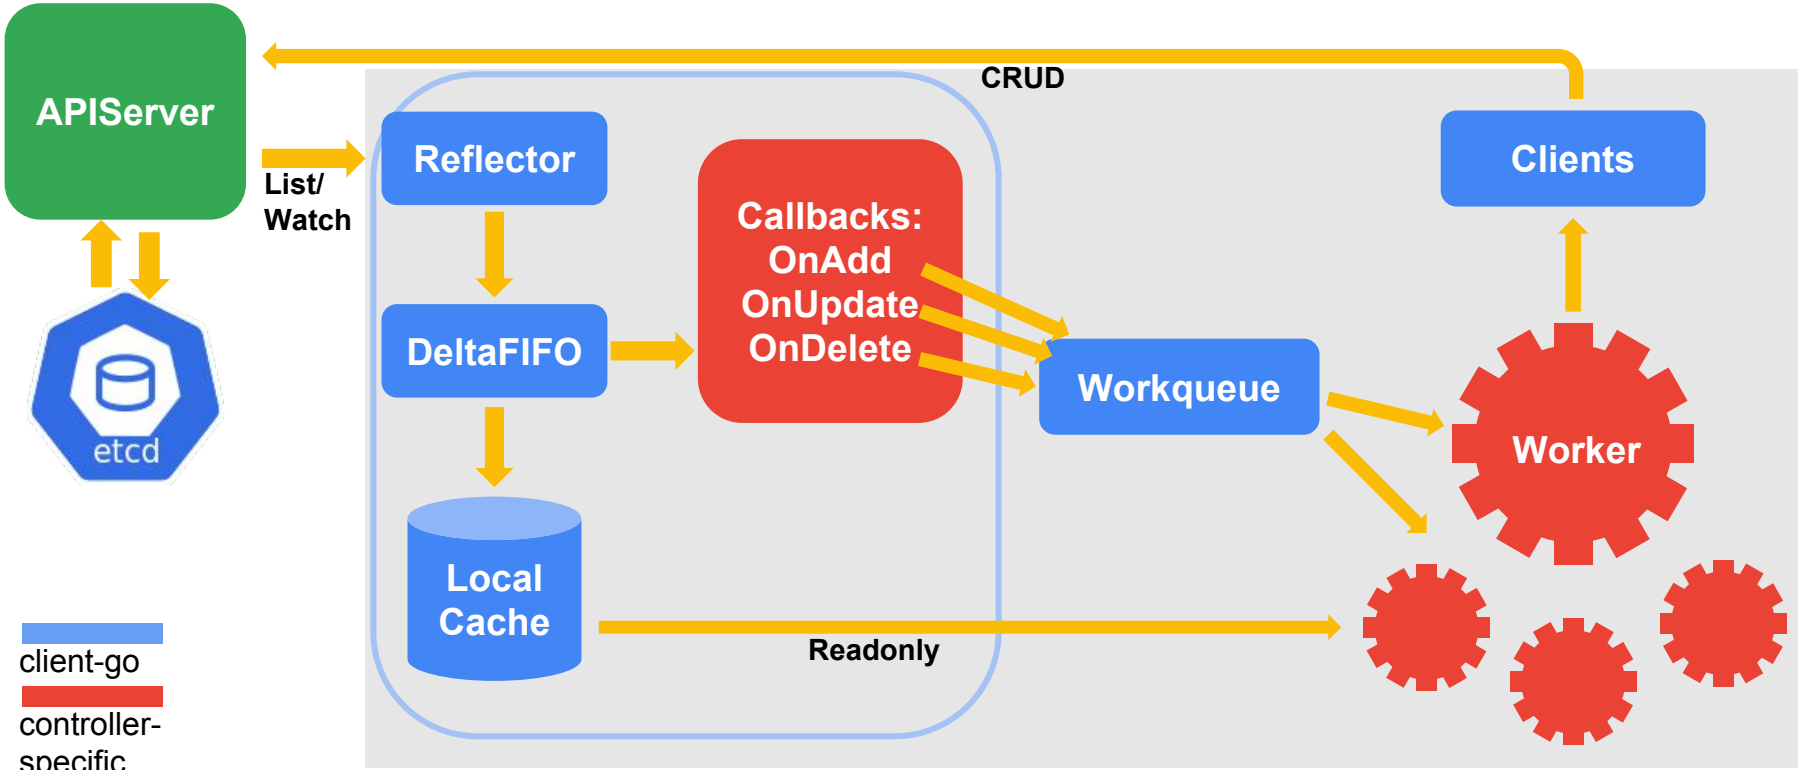

Kubernetes controller workflow

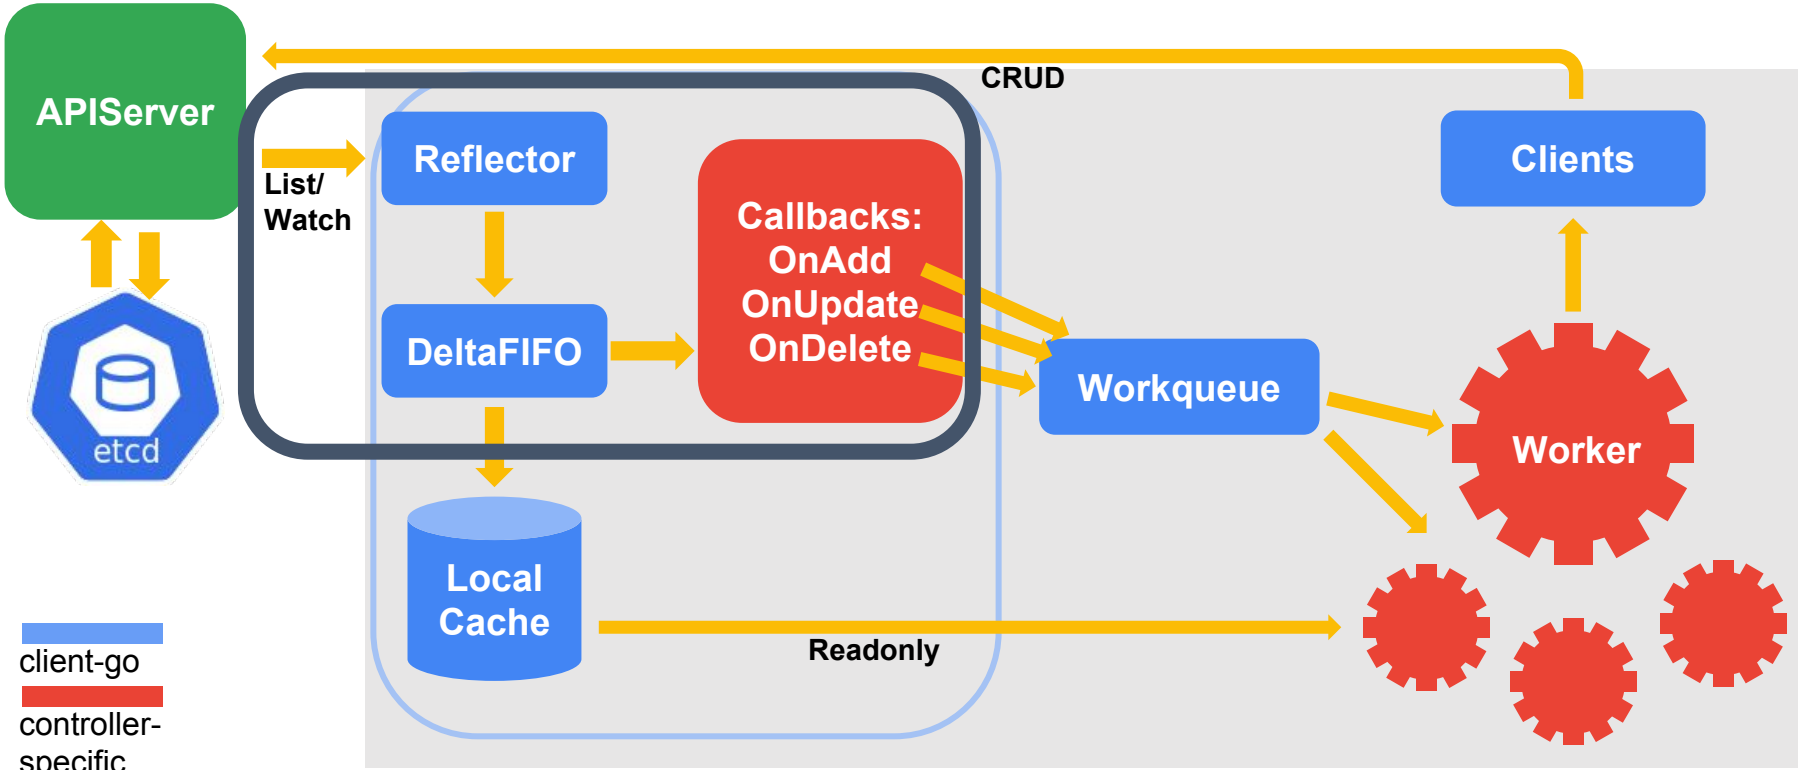

List and Watch


```
func ListAndWatch() error {
        // Explicitly set "0" as resource version - it's fine for the List()
        // to be served from cache and potentially be delayed relative to
        // etcd contents. Reflector framework will catch up via Watch() eventually.
        list, err = listWatcher.List(ListOptions{ResourceVersion: "0"})

        // Extract the actual ResourceVersion that we've listed
        resourceVersion = list.GetResourceVersion()

        for {
                options = ListOptions{ResourceVersion: resourceVersion}
                w, err = listWatcher.Watch(options)
                if err {
                        // most likely apiserver is not responsive. It doesn't make
                        // sense to re-list all objects because most likely we will
                        // be able to restart watch where we ended.
                        if err.IsError("connection refused") {
                                sleep(time.Second)
                                continue
                        }
                        // Watch closed normally (EOF) or unexpected error
                        HandleError(err)
                        return nil
                }
                // watchHandler watches w and keeps *resourceVersion up to date.
                watchHandler(w, &resourceVersion)
        }
}
```

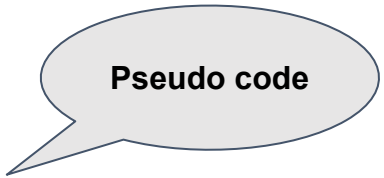

Recap: resourceVersion

Everything has a ResourceVersion:

- Individual API object (e.g. a Pod) has ResourceVersion when you GET one
- For a list of API objects (e.g. a PodList)
  - The entire list has a ResourceVersion
  - Each API object in items has ResourceVersion

The ResourceVersion of the top-level list is what should be used when starting a watch to observe events occurring after that list was populated.

Recap: resourceVersion

- ListOption in List Request
  - Unspecified: etcd
  - RV=0: APIServer cache
  - RV>0: the result is at least as fresh as given RV

Recap: resourceVersion

- ListOption in Watch Request
  - Unspecified: random time point
  - RV=0: the result is an "ADDED" event for every existing object followed by events for changes that occur after the watch was established
    - (main reason: backwards compatibility-- #13910)
  - Best practice: always specify last listed/watched RV

Watch Event on kube-scheduler, kube-controller-manager, kublet... (roycaihw)

Kubernetes controller workflow

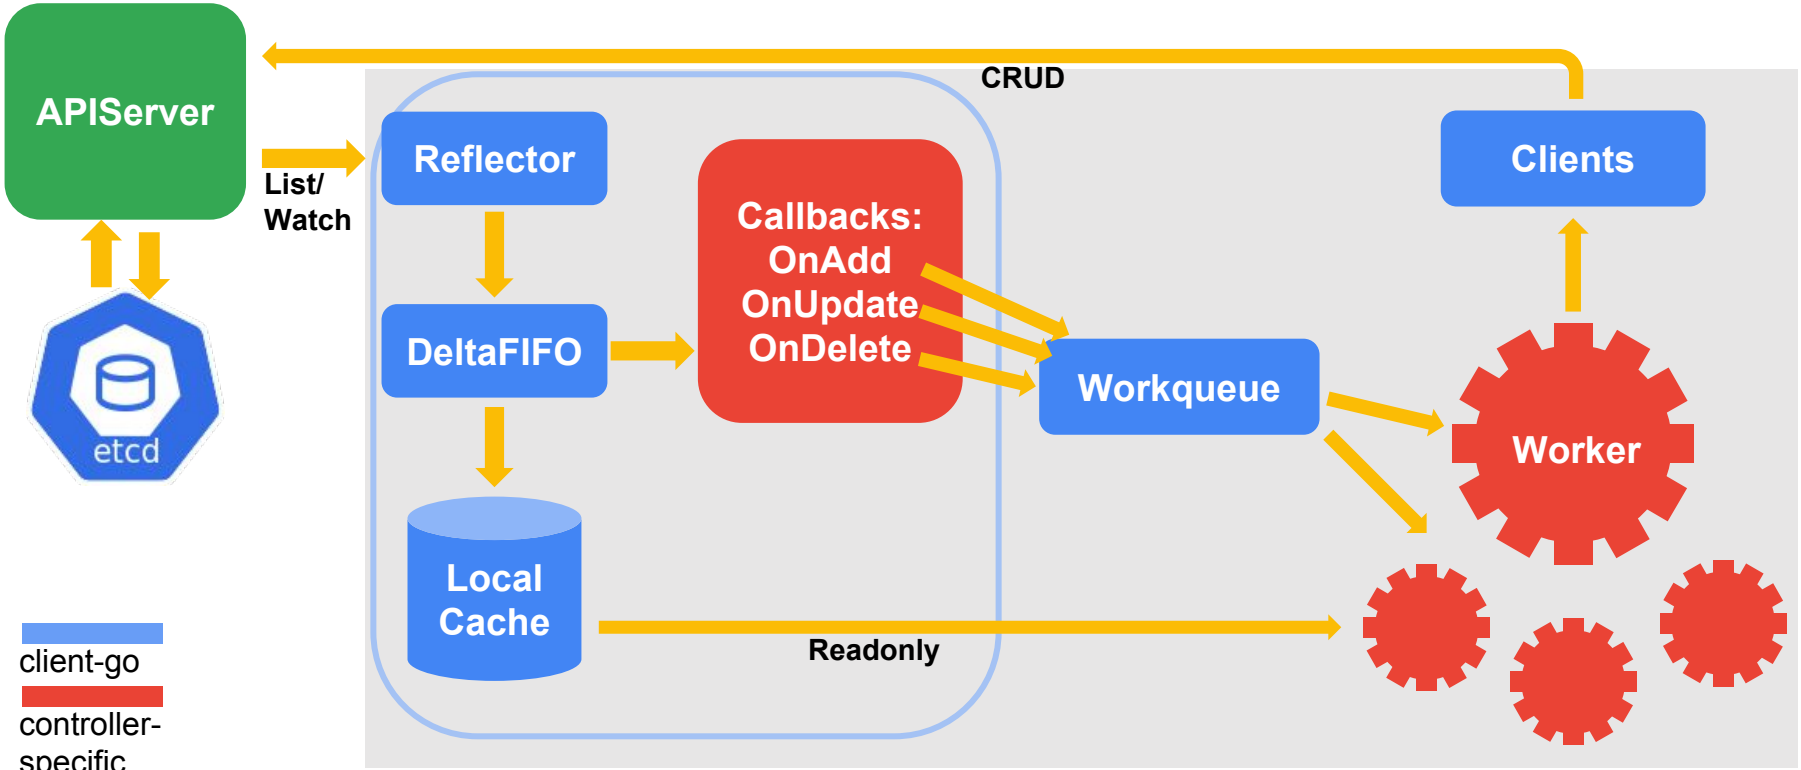

Scheduler


```
136  // New returns a Scheduler
137  func New(client clientset.Interface,
138         nodeInformer coreinformers.NodeInformer,
139         podInformer coreinformers.PodInformer,
140         pvInformer coreinformers.PersistentVolumeInformer,
141         pvcInformer coreinformers.PersistentVolumeClaimInformer,
142         replicationControllerInformer coreinformers.ReplicationControllerInformer,
143         replicaSetInformer appsinformers.ReplicaSetInformer,
144         statefulSetInformer appsinformers.StatefulSetInformer,
145         serviceInformer coreinformers.ServiceInformer,
146         pdbInformer policyinformers.PodDisruptionBudgetInformer,
147         storageClassInformer storageinformers.StorageClassInformer,
148         recorder record.EventRecorder,
149         schedulerAlgorithmSource kubeschedulerconfig.SchedulerAlgorithmSource,
150         opts ...func(o *schedulerOptions)) (*Scheduler, error) {
151
```

pkg/scheduler/scheduler.go

Watches:

- Node
- Pod
- PV
- PVC
- RC
- RS
- Stateful set
- Service
- PDB
- Storage class

Scheduler “business” logic


```
func InitEventHandlers(PodInformer) {
        // scheduled pod cache
        PodInformer.AddEventHandler(
                FilteringResourceEventHandler{
                        FilterFunc: func(obj interface{}) bool {
                                return IsScheduled(obj.Pod())
                        }
                        Handler: {
                                OnAdd: addPodToCache,
                                OnUpdate: updatePodInCache,
                                OnDelete: deletePodFromCache,
                        }
                }
        )

        // unscheduled pod queue
        PodInformer.AddEventHandler(
                FilteringResourceEventHandler{
                        FilterFunc: func(obj interface{}) bool {
                                return !IsScheduled(obj.Pod())
                        }
                        Handler: {
                                OnAdd: addPodToSchedulingQueue,
                                OnUpdate: updatePodInSchedulingQueue,
                                OnDelete: deletePodFromSchedulingQueue,
                        }
                }
        )
}
// SchedulingQueue is an interface for a queue to store pods waiting to be scheduled.
// The interface follows a pattern similar to cache.FIFO and cache.Heap and
// makes it easy to use those data structures as a SchedulingQueue.
type SchedulingQueue interface {
```

Pseudo code

Scheduler keeps in-memory:

- PodCache for scheduled pods
- SchedulingQueue for pods waiting to be scheduled

Having knowledge of the state of world enables scheduler to

- calculate resource usage on nodes
- be aware of available PVs that can be binded with pods
- schedule pods onto nodes

Watch enables on-demand information propagation, and avoids heavy information-polling loads.

Key Takeaways

Use Watch for your controller!

- It’s trustworthy!
- It’s efficient!


KubeCon
CloudNativeCon
North America 2018
?