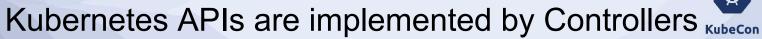
Unless otherwise specified...



```
Copyright 2017 The Kubernetes Authors.
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at
   http://www.apache.org/licenses/LICENSE-2.0
Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License.
```









- Respond to changes to Resources made by users
- Respond to changes to Cluster State

Controllers Are Simple



- ✓ watch resources,
- ✓ do some business logic,
- ✓ reconcile differences,
- ... and a bit of plumbing

How complicated could it be?

Sample Controller at a Glance



```
~/go/src/k8s.io/sample-controller $ wc -l < controller.go
429
~/go/src/k8s.io/sample-controller $ tree -I vendor
<snip>
28 directories, 52 files
```

Sample

```
re is "Controller: Run(threathness int, stopOn « cham atruct()) error (
  defer nurties NameLaCreek!)
   defer c.vertepaso, ShetDeard)
  aleg. Defer "waiting for inferior caches to synchri-
       primalt Metilic, runbisher, time, Second, scopthil
  glag beforetarted semants
  gleg. 2vfel"Shutting Some berkers")
use is "Controller! processionthorkItem! | book |
  ob1, stutdown := x.serkesses.6et()
  mer in Telepools into choosing across &
       var at test
```





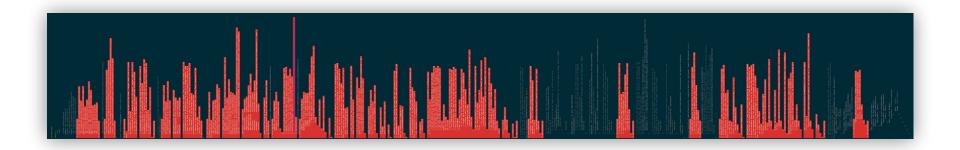
Sample Controller





This is not a flame chart





Note: Boilerplate code to copy-paste highlighted in Red

At least the comments are thorough...



```
// We call Done here so the workqueue knows we have finished
// processing this item. We also must remember to call Forget if we
// do not want this work item being re-queued. For example, we do
// not call Forget if a transient error occurs, instead the item is
// put back on the workqueue and attempted again after a back-off
// period.
// We expect strings to come off the workqueue. These are of the
// form namespace/name. We do this as the delayed nature of the
// workqueue means the items in the informer cache may actually be
// more up to date that when the item was initially put onto the
// workqueue.
// As the item in the workqueue is actually invalid, we call
// Forget here else we'd go into a loop of attempting to
// process a work item that is invalid.
```

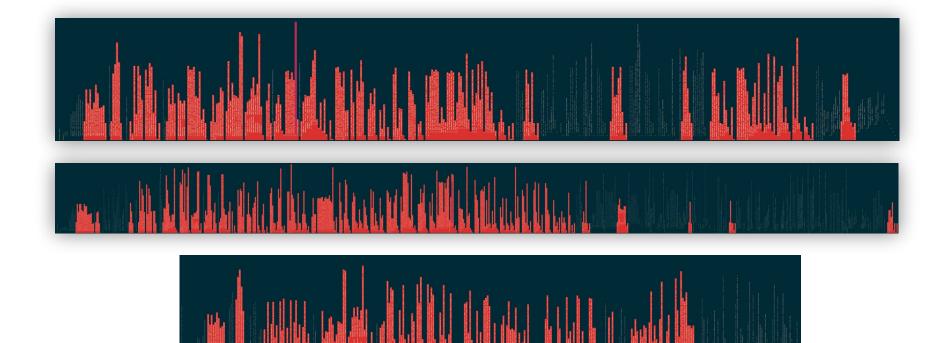
Controllers are Simple?

- Copy Clients to Reconciler Struct
- Create a Queue for Reconcile Requests (namespace/name)
- Add EventHandlers for all Object Types that you are interested in
 - Don't Enqueue for Deletion Events on the Reconciled Type
 - Unless you have a finalizer, then you should
 - Do Enqueue for Deletion Events on Owned Types
 - But find the owner of the object to Reconcile instead
 - Write Handlers to walk owners References to find parents
 - Be sure to handle Tombstones correctly
- Use Rateliniting when Requesting Forget the Object from the Queue Unless there is a transient error then don't Forger Actually maybe we don't need to do that anymore?
 - Write Logic to Turn Objects Into Keys
 - Type Cast Request to String
 - Error Handling if this fails
- Start Worker Threads
 - But not too many
- Start Informers
- Wait For Cache To Sync ge State

Namespace | Name Split String Into Check for unowned Objects Using number

Which is the real Sample Controller?





Sample Controler + Core Kubernetes Controller + Third Party Operator

Controllers Should Be Simple



65% of the sample controller code is...

- Unmodified Copy-Pasted Functions
- Initialization and Instantiation of Deps
- Plumbing and Wiring

Native Controller Support in Go 1.12!



controller-runtime (proper noun)

a Go standard library for controllers

```
package main
import (
     "controllers"
)
```

Just Kidding

Seriously though, use this instead



controller-runtime (proper noun)

a Go standard library for controllers

```
package main
import (
        "sigs.k8s.io/controller-runtime/pkg/builder"
)
```

The magic of abstractions...



Patterns

- Watch Objects to Reconcile
- Watch Generated Objects
- Watch Objects that Map to Reconciled Objects

Utilities

- Pre-Split Namespace/Name
- Universal Client (no more wiring generated clients)
- Helper for setting Owners References
- Helper for Upserting Objects

Create Controller



```
type Controller struct {}
func Start() {
   ctrl, _ := builder.SimpleController().Build(&Controller{})
   ctrl.Start(signals.SetupSignalHandler())
}
```

Watch Object





```
klog.Fatalf("Error building example clientset: %s", err.Error())
      AddFunc: controller.engueueFoo.
      UpdateFunc: func(old. new interface()) {
func (c *Controller) enqueueFoo(obj interface{}) {
  var key string
      runtime.HandleError(err)
```

```
type Controller struct {
   Client client Client
func (ct *Controller) InjectClient(c client.Client) error {
   ct.Client = c
    return nil
func Start() {
   ctrl, := builder.SimpleController().
        ForType(&v1alpha1.Foo{}).
        Build(&Controller{})
    ctrl.Start(signals.SetupSignalHandler())
```

Watch Owned Objects



```
type Controller struct {
   Client client.Client
func (ct *Controller) InjectClient(c client.Client) error {
   ct.Client = c
   return nil
func Start() {
   ctrl, := builder.SimpleController().
       ForType(&vlalpha1.Foo{}).
       Owns(&appsv1.Deploment{}).
       Build(&Controller{})
   ctrl.Start(signals.SetupSignalHandler())
```

Sophisticated Watch APIs



```
// .ForType(&corev1.Pod{})
c.Watch(&source.Kind{Type: &corev1.Pod{}},
    &handler.EnqueueRequestForObject{})
```

The. For Type builder function is really just a call to .Watch with EnqueueRequestForObject

Sophisticated Watch APIs



```
// .Owns(&appsv1.ReplicaSet{})
c.Watch(&source.Kind{Type: &corev1.Pod{}},
    &handler.EnqueueRequestForOwner{
    IsController: true, OwnerType: &appsv1.ReplicaSet{}})
```

The. Owns builder function is really just a call to .Watch with EnqueueRequestForOwner

Sophisticated Watch APIs



```
// Not Present In Builder - Map One Object to other Objects
c.Watch(&source.Kind{Type: &appsv1.Deployment{}},
   &handler.EnqueueRequestsFromMapFunc{ToRequests: mapFn})
```

Users can implement their own patterns by providing a function to map events to objects

Extending Watch APIs





Watch(src source.Source, eventhandler handler.EventHandler, predicates ...predicate.Predicate) error type EventHandler interface { Create(event.CreateEvent, workqueue.RateLimitingInterface) Update(event.UpdateEvent, workqueue.RateLimitingInterface) Delete(event.DeleteEvent, workqueue.RateLimitingInterface) Generic(event.GenericEvent, workqueue.RateLimitingInterface) type Source interface { Start(handler.EventHandler, workqueue.RateLimitingInterface, ...predicate.Predicate) error

Shared Dependencies





```
// Implement these to have dependencies Injected
type Client interface {
    InjectClient(client.Client) error
type Config interface {
    InjectConfig(*rest.Config) error
type Scheme interface {
    InjectScheme(scheme *runtime.Scheme) error,
type Stoppable interface {
    InjectStopChannel(<-chan struct{}) error</pre>
type Injector interface {
    InjectFunc(f Func) error
```

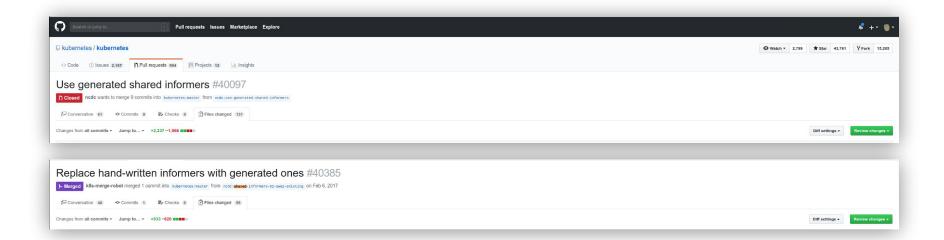
Reconcile Utilities



```
// Sample Controller
OwnerReferences: []metav1.OwnerReference{
    *metavl.NewControllerRef(foo, schema.GroupVersionKind{
             samplevlalphal.SchemeGroupVersion.Group,
    Group:
    Version: samplevlalphal.SchemeGroupVersion.Version,
    Kind:
            "Foo"})}
// Controller-Runtime
controllerutil.SetControllerReference(foo, deployment, c.scheme)
```

For Developers...





For Ops...



Standard...

- Behavior
- Logging
- Metrics
- Error handling and degradation
- Maintenance

For Platform Developers...

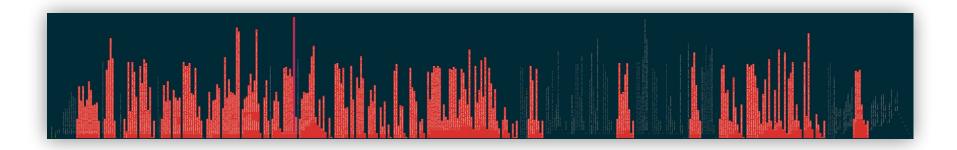


Build high-level abstractions with common building blocks:

- Kubebuilder
- Operator SDK
- Maestro Declarative Operator (via kubebuilder)
- AddOnOperator (via kubebuilder)

Remember this...





It fits!





In conclusion...



Controllers Are Simple

```
package main
import (
     "sigs.k8s.io/controller-runtime/pkg/builder"
)
```

Applause...

